

Creating Service Clouds with Kubernetes, CoreOS, Containers and Co

thomas@endocode.com

ENDOCODE

CONTAINER

CONTAINER

- container
 - resources
 - access rights
 - images
- formats
 - rkt
 - aci
- registries

CONTAINER VS VIRTUALIZATION

Topic	Container	Virtualisation	
Isolation	OS Level, OS namespaces	CPU Level: Ring 0/Ring 3	
foreign CPU	no	yes, with emulation	
foreign kernels, OS	no	yes	kernel is common
emulated devices	no	yes	security
host devices	direct	virtio driver	security
CPU performance	100%	95%	
IO performance	100%	<<100%	
root isolation	yes	yes	USER directive
CPU cache attacks	easy	possible	PoC ?

KUBERNETES

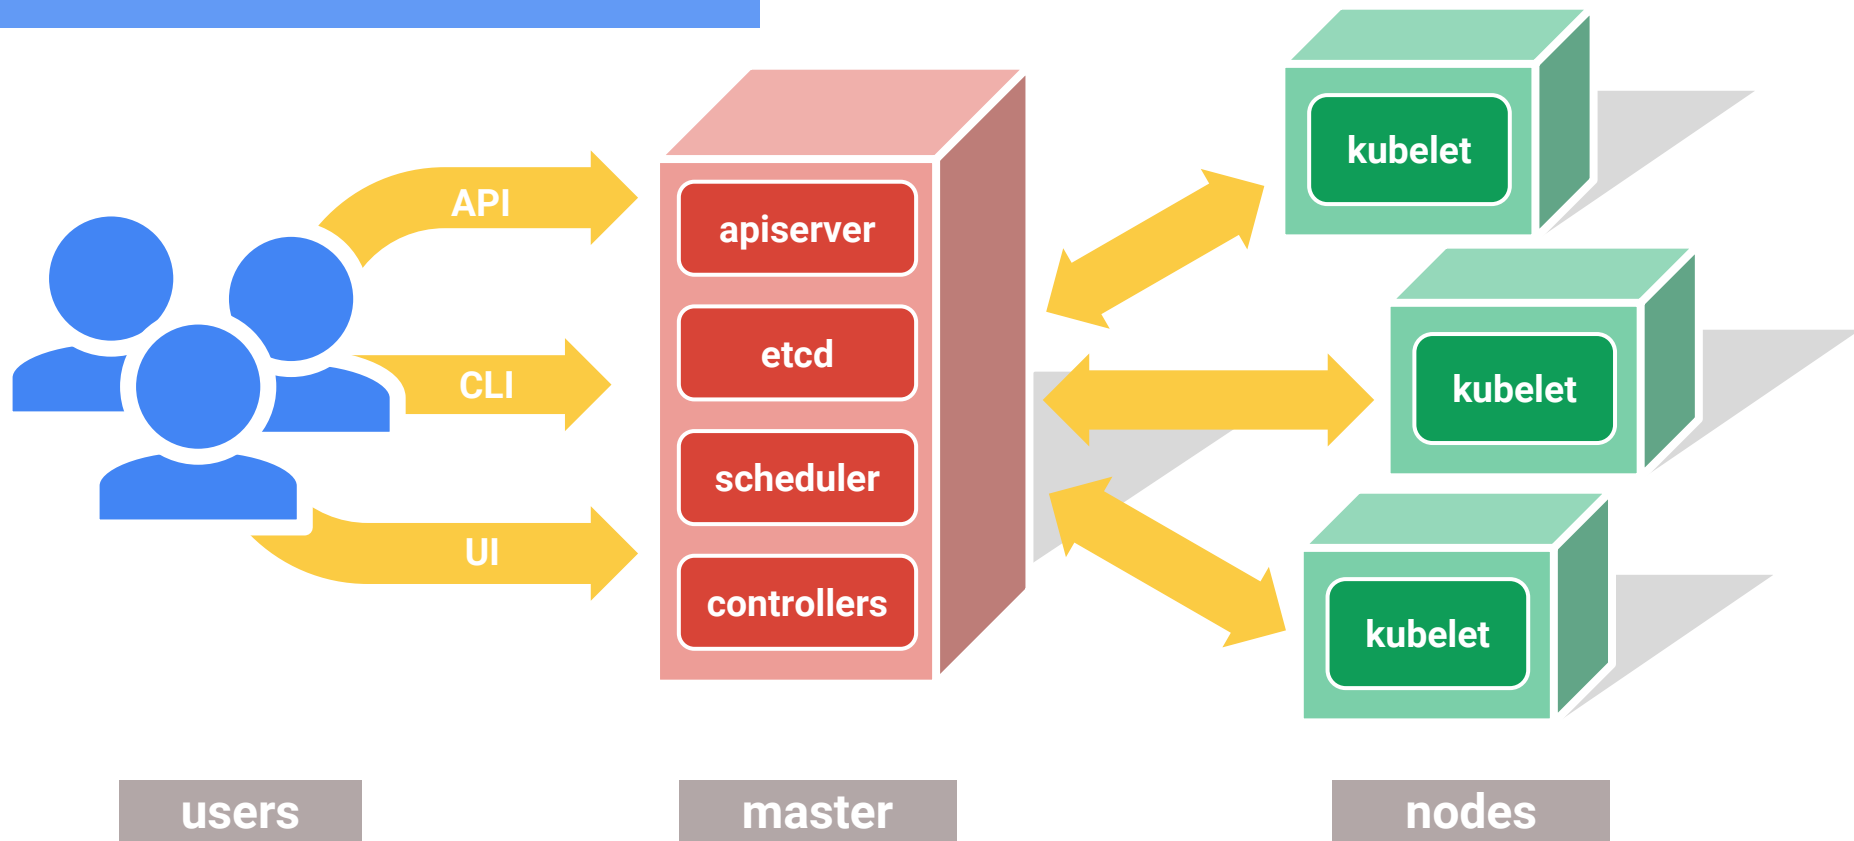
Greek for *“Helmsman”*; also the root of the words *“governor”* and *“cybernetic”*

- Runs and manages containers
- Inspired and informed by Google’s experiences and internal systems
- Supports multiple cloud and bare-metal environments
- Supports multiple container runtimes
- **100% Open source**, written in Go


Manage applications, not machines



The 10000 foot view

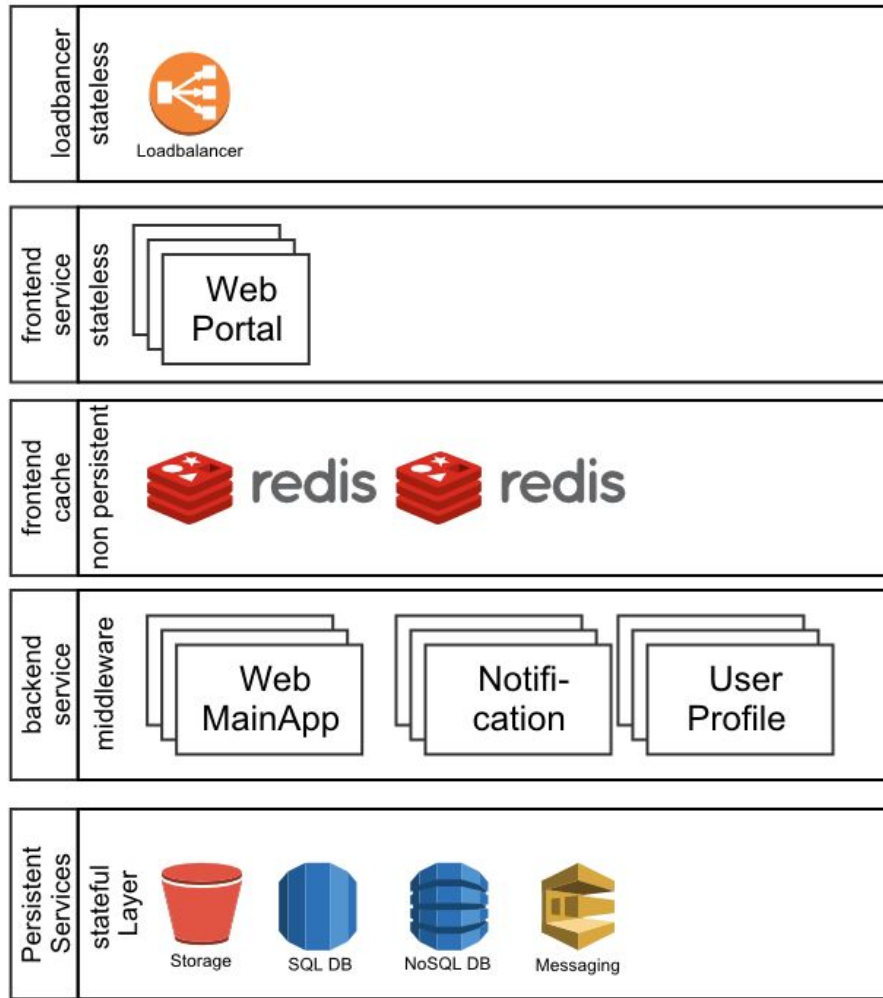


All you really care about



ARCHITECTURE

- strict layered architecture
 - separation of stateless services “cattle”
 - and persistent data “pets”
- inside the pods
 - developers are free to use what they want
 - contract is binding to the outside



Kubernetes Building Blocks

- Pods
- Replication Controller
- Services
- Persistence
 - External Ressources
 - Databases
 - Local Disk
 - Cloud Storage
 - Replication
 - Backup

STATELESS AND STATEFUL SERVICES

- where to keep state? A trade-off
 - provider → lock-in
 - self-managed → overhead
- cattle, no pets
- mindset: ephemeral deployment units

FROM VMs TO PODS

OS instances  microservices in Pods

- pods are containers sharing the same fate
 - created together
 - running on same node
 - terminating together
 - one network address
 - shared volumes

FROM VMs TO PODS

VM cluster  Pods running on Kubernetes

- cattle: stateless containers
- pets: databases

configuration management  separation of build time and run time

GETTING STARTED WITH KUBERNETES

PREREQUISITES

- Networkprovider
- Storage
- Loadbalancer Provider
- Ops Use Cases
 - add node
 - delete node
 - replace failing node
 - update kubernetes
- installation
 - cloud provider
 - bare metal

PUBLIC VS PRIVATE CLOUD

- GKE
 - ready to rock
 - scales very well
- AWS
 - need to set up the Ops cases
 - rich set of services
 - less need for pets
 - replicated databases
- Your Private Cloud
 - loadbalancer providers
 - implement
 - manual
 - lots of storage providers
 - nfs
 - ceph
 - bare metal
 - local disks
 - SSDs

LAST OPS TASKS

- Backups
 - replication
 - database
 - storage
 - snapshots
 -
- Monitoring
 - basic
 - business KPIs
 - prometheus
 - graphite
 -
- Alerting
 - based on monitoring
 -